

Jane Wu

(412)-584-1966 | wujane001@gmail.com | <https://www.linkedin.com/in/janewu-zichu-wu/> | <https://jjane-wu.github.io/>

EDUCATION

Carnegie Mellon University

Aug 2024 - Dec 2026

Master of Information Systems Management (GPA: 3.9/4.0)

Pittsburgh, PA

- Coursework: A/B Testing (Teaching Assistant), Machine Learning, Statistical Analysis, Agentic AI, Product Management

New York University

Sep 2020 - May 2024

Bachelor of Science, Business and Finance, Minor in Mathematics (GPA: 3.8/4.0)

New York, NY

WORK EXPERIENCE

United Nations Joint Staff Pension Fund

June 2026 - Sep 2026

Data Scientist Intern

New York, NY

- Owned a **multi-agent due-diligence assistant** on **AWS Bedrock and Kubernetes (EKS)**, enabling the private-equity team to self-serve validated return metrics (MOIC, IRR, loss ratio) from fund managers' any-format track records in **~1 minute** instead of hours
- Implemented its **LLM tool-calling workflow** over schema-inference, validation and return-analytics tools on a **PostgreSQL (pgvector)** backend, guarded by confidence gates, **15 validation rules** and analyst review; scoped PRD with non-technical investment analysts

TikTok

May 2026 - Aug 2026

Product Data Scientist Intern

Shanghai, China

- Optimized a **LightGBM** model of how creator coverage drives livestream revenue to guide creator allocation, flagging under-promoted SKUs that grew **revenue 10-29%** the next campaign
- Led the readout of a creator-randomized **A/B test** of a CRM ranking algorithm (CUPED-adjusted metrics): a **pre-AA check** traced a spurious **~5%** revenue drop to a contaminated split, and funnel analysis pinpointed a **~10%** fall in creator adoption as the next fix
- Built the data pipeline for an internal **attribution-analysis AI agent** that runs the team's attribution model and writes category analyses for operations, designed to cut turnaround from **3-5 days to 1 hour**
- Presented **5 market research reports** to leadership, benchmarking Southeast Asia livestream against Douyin for category strategy

Xiaomi

May 2025 - Nov 2025

Data Scientist Intern

Beijing, China

- Fine-tuned Qwen2.5-14B with **LoRA**, chosen over 7B and 72B on the accuracy-cost trade-off, and fixed out-of-memory errors by cutting **GPU memory ~75%** (INT8 quantization, gradient checkpointing), reaching **84% accuracy** across 42 auto-insurance failure reasons
- Added a **LangChain RAG** pipeline (ChromaDB, few-shot examples) that lifted long-tail macro-F1 **12.5%** and the full system to **0.88 F1** (57% over BERT), routing low-confidence labels to human review
- Productionized **Spark** pipelines with Airflow to process 2+ TB of daily user-behavior logs in Iceberg, powering funnel-monitoring dashboards and downstream credit risk model

LianHai Capital Asset Management

Jun 2024 - Aug 2024

Sector Data Scientist Intern

Beijing, China

- Built a **factor back-testing framework** on ~5,000 A-shares (2015-2024) using daily **constrained cross-sectional regressions** for industry baselines, decile sorts and IC/IR, and automated a daily 9-factor review for portfolio managers

PROJECTS

Movie Recommendation System in Production | <https://github.com/cmu-seai/f26-project-31-the-silence-of-the-llms>

- Deployed a real-time recommendation API on live **Kafka** traffic with fallback routing that answers every request, at **0.21 ms p95** model latency (600 ms budget)
- Built **CI/CD** (GitHub Actions, **Docker**, 56 tests per PR, automated retraining) and deployed **Prometheus and Grafana** monitoring
- Designed a **cold-start model** for the 41% of users with no history; an online A/B test raised the 30-minute **watch rate by 6%** ($p < 0.05$)

Latent Diffusion Model with Classifier-Free Guidance | https://github.com/Jackiebibili/idl_diffusion

- Built a latent diffusion pipeline (VAE, U-Net, **DDIM**, classifier-free guidance) on ImageNet-100, trained on CMU HPC

Face Recognition and Verification with CNNs | PyTorch, ResNet-18 from scratch

- Implemented **ResNet-18 from scratch** in PyTorch and trained it on 8,631 identities with mixed-precision training and augmentation
- Verified faces via cosine similarity of L2-normalized 512-d embeddings, reaching **3.79% EER**; tuned with cosine-annealing restarts

SKILLS

- **Languages & Frameworks:** Python, SQL, R, Java, Scala, Bash, TypeScript, JavaScript, React, FastAPI, REST APIs
- **AI & ML:** PyTorch, TensorFlow, Scikit-Learn, Hugging Face, LangChain, NLP, AI Agents, Prompt Engineering
- **Data & Cloud:** Spark, Hive, Hadoop, Airflow, Kafka, Flink, Snowflake, Databricks, DBT, AWS, GCP, Docker, Kubernetes
- **Analytics & Tools:** A/B Testing, Causal Inference, Time Series Forecasting, Tableau, Power BI, Looker, Excel, Figma, Git, Agile